

Building Autonomous AI Agents: A Comprehensive Primer

October 7, 2026

Table of Contents

1	Introduction	2
1.1	What This Covers	2
1.2	What This Is Not	3
1.3	The Default Production Path	3
2	Foundations and Agent Boundaries	4
2.1	What Is an Autonomous AI Agent?	4
2.2	The Core Components	4
3	The Model and Tool Boundary	6
3.1	The Misconception: LLMs as Black Boxes	6
3.2	The Trick: Tool Use and Structured Outputs	6
3.3	Native Tool Use APIs	7
3.4	The Tool Boundary Is a Trust Boundary	8
3.5	Why This Works	8
4	The Production Execution Loop	9
4.1	The Agent Lifecycle	9
4.2	From a Loop to a Durable Workflow	11
4.3	Concrete Example: A Customer Support Agent	11
5	Tools and Actions	13
5.1	What Makes a Good Tool?	13
5.2	Categories of Tools	15
5.3	Tool Orchestration Patterns	15
6	Memory and Context	16
6.1	The Context Window Problem	16
6.2	Memory Hierarchy	16
6.3	Managing Memory in Practice	17
7	Reference branch: Alternative Patterns, Architectures, and Frameworks	18

7.1	ReAct Pattern (Reasoning + Acting)	18
7.2	Planning Patterns	19
7.3	Error Handling and Recovery	19
7.4	The Execution Loop Implementation	19
7.5	Frameworks, SDKs, and Runtimes (September 2026 Snapshot)	20
8	Production Risks and Failure Modes	22
8.1	Challenge 1: Hallucination	22
8.2	Challenge 2: Context Length Limits	23
8.3	Challenge 3: Cost and Latency	23
8.4	Challenge 4: Ensuring Reliability	24
8.5	Challenge 5: Debugging and Understanding	25
9	Build the Production Path	25
9.1	Start Simple	25
9.2	Then Gradually Add Complexity	26
9.3	Testing an Agent	26
9.4	Path step 1: Define Scope and Guardrails	27
9.5	Path step 2: Define Success Metrics Before Coding	27
9.6	Path step 3: Design Tool Contracts	27
9.7	Path step 4: Write the System Prompt Like an Operating Manual	28
9.8	Path step 5: Implement the Execution Loop with Hard Stops	28
9.9	Path step 6: Build a Realistic Test Suite	28
9.10	Path step 7: Add Human-in-the-Loop for Sensitive Paths	28
9.11	Path step 8: Roll Out Gradually	29
9.12	Path step 9: Operate, Observe, Iterate	29
9.13	What “Working” Looks Like	29
10	Conclusion	29
10.1	Current Limitations (September 2026)	29
10.2	Reference branch: Emerging Patterns	30
11	Further Reading and Resources	31

1 Introduction

1.1 What This Covers

This primer covers how autonomous AI agents work at a technical level: the execution loop, tool use, memory management, multi-agent patterns, and common failure modes. It walks through design decisions you will face when building agents that act reliably in the real world, and includes a worked end-to-end example.

1.2 What This Is Not

This is not a guide to training or fine-tuning LLMs. It is not a benchmark comparison of AI providers, and it does not cover prompt engineering in isolation from agentic execution. If you want a survey of agent frameworks without implementation depth, this is probably not the right starting point — this primer assumes you intend to build something.

Freshness note (October 7, 2026): The terminology and editorial review were updated on this date; the framework and product snapshot remains September 2026. Framework capabilities, provider SDK features, model pricing, and benchmark rankings move monthly. Treat product and protocol references as dated guidance and re-check the linked official docs before implementation.

1.3 The Default Production Path

Use this primer as one end-to-end path when you are building an agent for a real workflow. Read the numbered sections in order: define the agent and its trust boundary, understand the model/tool decision boundary, design the execution loop, define narrow tools, manage state and memory, account for failure modes, and then implement and operate the production path. The customer-support walkthrough in the execution-loop section is the running example; the second support walkthrough in the build section turns that example into an implementation plan.

The section titled **Reference branch** contains useful alternatives—ReAct, planning variants, multi-agent coordination, provider SDKs, MCP, and A2A—but those are branches from the default single-agent, tool-grounded workflow. Choose one only when its additional state, latency, trust boundary, or operational complexity solves a problem that the default path cannot.

When you hear about AI agents doing autonomous work—making API calls, retrieving data, making decisions, and executing complex workflows—you’re likely hearing about a fascinating but often misunderstood technology. The core confusion typically stems from a simple question: *if AI systems like large language models (LLMs) are trained to generate text, how do they become agents that actually do things?*

The answer is elegantly simple: **LLMs don’t become agents by themselves. Instead, they become the reasoning engine inside a larger system that**

combines their decision-making capabilities with tools, memory, and structured processes. This primer explains how that system works, why it's powerful, and what you need to understand to build one.

If you're coming from a software engineering background, a useful mental model is this: an agent is just a program with an LLM in the decision loop. If you're coming from product or operations, think of it as a very capable junior operator that can read instructions, use software tools, and report what it did.

We'll keep this practical and conversational. The goal here is not to bury you in buzzwords. The goal is to help you build something that still works when inputs are messy, APIs are flaky, and users are unpredictable.

2 Foundations and Agent Boundaries

2.1 What Is an Autonomous AI Agent?

An autonomous AI agent is a software system that:


1. **Perceives** its environment (accepts input, reads data, monitors state)
2. **Reasons** about what to do next (using an LLM)
3. **Acts** to accomplish goals (calls APIs, executes functions, makes changes)
4. **Adapts** based on feedback and iterates — through memory and updated context, not through retraining; the model's weights don't change mid-task

The crucial distinction: An agent is not just an LLM chatbot. A chatbot waits for human input and generates text in response. An **agent takes initiative**—it decides what actions to take without waiting for the human to specify every step.

A quick litmus test: if the system can only respond with text, it's a chat assistant. If it can choose and execute actions toward a goal, it's agentic.

2.2 The Core Components

Every autonomous AI agent has these essential parts:



1. The LLM (Large Language Model) The “brain” of your agent. Models like Claude, GPT, Gemini, Llama, etc. are exceptionally good at understanding complex instructions and generating structured reasoning. They consume input text and produce output—including decisions about what actions to take.

2. Tools / Actions These are the agent’s hands. Tools are functions the agent can call: making HTTP requests, querying databases, reading files, sending emails, etc. The agent decides *which* tools to use based on its reasoning.

3. Memory & Context The agent needs to maintain state: conversation history, facts about the current task, previous decisions, learned information. This prevents the agent from forgetting what it’s done and why, and helps it make coherent decisions.

4. Execution Loop The orchestration layer that repeatedly: asks the LLM what to do next → executes that action → feeds results back to the LLM → repeats until the goal is achieved or a stopping condition is met.

Mini example: an IT access agent receives “Grant dashboard access.” It checks user identity, checks role policy, files an access request, and sends a confirmation. That simple loop is already a complete autonomous workflow.

3 The Model and Tool Boundary

3.1 The Misconception: LLMs as Black Boxes

Many people assume LLMs are mysterious, but for agent-building purposes, you can think of them simply as:

Pattern matchers that are excellent at understanding context and producing structured outputs.

When you train an LLM on vast amounts of text—code, documentation, reasoning traces, instructions—it learns patterns. One of those patterns is: “When someone asks me to accomplish a goal, here’s how reasoning should flow.”

Another way to think about this: the model is very good at choosing the next best action given context. In agent systems, that action is often a tool call. That is why tool design and prompt clarity matter so much.

3.2 The Trick: Tool Use and Structured Outputs

Here’s where the magic happens. Instead of just asking an LLM to “solve this problem,” you structure the interaction so the LLM outputs **instructions for what to do next**.

The example below shows the underlying idea using plain text instructions, so you can see exactly what’s happening with nothing hidden behind an API. It’s conceptually accurate but it is not how you’d build this today — the subsection right after the example explains what production systems actually do.

Example 1: Simple tool use

System prompt: “You are an agent. You can call these tools:

- search(query): searches the web
- read_file(path): reads a file
- write_file(path, content): writes a file

When you want to use a tool, respond in this format:

```
<tool_call>
{"name": "tool_name", "arguments": {...}}
```

```
</tool_call>
"
```

User: "Find the current Bitcoin price and save it to bitcoin_price.txt"

LLM output:

"I'll search for the current Bitcoin price and save it for you.

```
<tool_call>
```

```
{"name": "search", "arguments": {"query": "current Bitcoin price"}}
```

```
</tool_call>"
```

The system then:

1. Parses the LLM's output to extract the tool call
2. Executes `search("current Bitcoin price")`
3. Gets results like "Bitcoin: \$42,500"
4. Feeds those results back: "Search results: Bitcoin: \$42,500"
5. Asks the LLM what to do next

This loop repeats until the LLM says the task is complete.

3.3 Native Tool Use APIs

The tag-based example above works because you told the model to imitate a format in the system prompt — nothing stops it from getting the JSON slightly wrong, wrapping it in markdown, or emitting it mid-sentence. You'd then own a regex/parsing layer responsible for handling every way a model can almost-but-not-quite follow instructions. This is roughly how agents were built in early 2023, and it is fragile in exactly the way you'd expect a text-parsing layer bolted onto a language model to be.

Every major provider now ships a native tool-use API instead. You pass your tool definitions as structured schemas in the API request itself — not as prose in the system prompt — and the model returns a typed `tool_use` content block: a proper object with a name and arguments, not text you have to extract them from. You execute the tool and send the result back as a `tool_result` block, and the loop continues from there. Anthropic, OpenAI, and Google all support this shape, and the same pattern now appears across a broader set of providers as well; the `tool_use/tool_result` terminology above is Anthropic's, but the concept — schemas in, typed call out — is now broadly convergent. Models are also specifically trained against this interface, so the reliability gain isn't just "no parsing bugs," it's "the model is better at using tools when they're presented this way."

The practical implication: skip the prompt-tag approach entirely unless you're working with a raw base model that has no tool-calling interface at all. For any current frontier model, define your tools through the API's native tools parameter and read `tool_use` blocks off the response. That's the standard mechanism, not an optimization layered on top of the standard mechanism — the pseudocode later in this primer assumes it.

Same pattern, different domain: an expense-audit agent can fetch a report, compare each line item to policy, request missing receipts, then approve or escalate. Same loop, different tools.

3.4 The Tool Boundary Is a Trust Boundary

The model proposes an action; application code decides whether that action is allowed and performs the side effect. Keep those responsibilities separate. Tool schemas constrain shape, but they do not provide authorization. A valid `delete_customer` call can still be the wrong action for the current user, tenant, policy, or workflow state.

Treat retrieved pages, emails, documents, and tool results as untrusted data. They may contain prompt-injection text that looks like an instruction. Put authorization, tenant checks, rate limits, idempotency keys, and approval gates in the tool or service layer, where they cannot be bypassed by a persuasive model response. For irreversible actions, pause before execution and show the human the exact arguments, target, and expected consequence.

The current ecosystem exposes this boundary through several converging interfaces: provider SDKs offer guardrails, handoffs, tracing, and sometimes sandboxed execution; [Model Context Protocol \(MCP\)](#) standardizes tool and resource connections; and [Agent2Agent \(A2A\)](#) addresses communication between independently deployed agents. These are integration layers, not replacements for application authorization.

3.5 Why This Works

LLMs have learned from millions of examples of:

- Reasoning traces (showing step-by-step thinking)
- Programming (where functions are called and their outputs used)
- Instructions (where tasks are broken into steps)


So when you frame a problem as “call this tool, use the result to decide what to do next,” the LLM naturally generates reasonable plans.

This also explains why loop-based agents often outperform one giant prompt. The model doesn’t need to predict the whole future in one shot. It can take one step, observe reality, and adapt.

4 The Production Execution Loop

4.1 The Agent Lifecycle

Every autonomous agent follows a repeating cycle:



In production, each iteration should be observable. Log the prompt summary, chosen action, tool arguments, result, and elapsed time. That one habit makes debugging dramatically easier.

4.2 From a Loop to a Durable Workflow

A short interactive loop can keep state in memory. A long-running or approval-gated agent needs durable state outside the process. Persist a workflow identifier, current step, tool-call status, model/prompt versions, and the minimum data needed to resume. Make side effects idempotent or record an idempotency key before retrying them; otherwise a crash after a successful payment or email can cause a duplicate when the agent resumes.

Checkpointing also changes the design of human review. The agent should pause with a reviewable action request, not merely ask the language model to “be careful.” The reviewer can approve, edit, or reject the exact call, and the workflow resumes from saved state. [LangGraph persistence](#) and [interrupts](#) are concrete examples of this pattern; the same separation can be implemented with a database-backed state machine or a durable workflow engine.

4.3 Concrete Example: A Customer Support Agent

Let’s walk through a real example to see this in action.

Setup:

- Task: “Handle customer complaint about order #12345: customer says ‘Widget A arrived broken.’”
- Policy: replace the item if it’s in stock; only refund if a replacement isn’t available.
- Available tools: `get_order()`, `check_inventory()`, `apply_refund()`, `ship_replacement()`, `send_email()`
- Goal: Resolve the issue in line with policy and confirm resolution with the customer

Iteration 1:

Prompt sent to LLM:

"You are a customer support agent. Policy: replace the item if it's in stock; only refund if a replacement isn't available.

Task: Handle customer complaint about order #12345: customer says 'Widget A arrived broken.'

Available tools: `get_order()`, `check_inventory()`, `apply_refund()`, `ship_replacement()`, `send_email()`

Previous context: None (first step)

What should you do first?"

LLM responds:

"I should first retrieve the order details to understand the issue.

<tool_call>

{"name": "get_order", "arguments": {"order_id": "12345"}}

</tool_call>"

System executes: Returns order details

```
{
  "id": "12345",
  "customer": "john@example.com",
  "items": [
    {"name": "Widget A", "price": 29.99},
    {"name": "Widget B", "price": 19.99}
  ],
  "status": "delivered",
  "delivery_date": "2026-05-15"
}
```

Iteration 2:

Prompt sent to LLM:

"Previous context:

- Retrieved order #12345: Widget A (\$29.99) and Widget B (\$19.99), delivered May 15
- Customer says: 'Widget A arrived broken.'

What should you do next?"

LLM responds:

"The customer received a broken Widget A. Per policy, I should check whether a replacement is available before considering a refund.

<tool_call>

{"name": "check_inventory", "arguments": {"item": "Widget A", "quantity": 1}}

</tool_call>"

System executes: Returns inventory status

```
{"item": "Widget A", "in_stock": true, "quantity": 15}
```

Iteration 3:

Prompt sent to LLM:

"Previous context:

- Order #12345: Widget A (\$29.99, broken) and Widget B (\$19.99), delivered May 15
- Widget A in stock: 15 units available
- Customer says: 'Widget A arrived broken.'

What should you do next?"

LLM responds:

"Widget A is in stock, and policy says to replace rather than refund when that's the case. I'll ship a replacement instead of issuing a refund.

<tool_call>

{"name": "ship_replacement", "arguments": {"order_id": "12345", "item": "Widget A"}}

</tool_call>"

System executes: Returns confirmation

```
{"success": true, "replacement_id": "RP-2026-0515-001"}
```

Iteration 4:

Prompt sent to LLM:

"Previous context:

- Order #12345: Widget A (broken, replacement shipped) and Widget B
- Replacement shipped: RP-2026-0515-001

What should you do next?"

LLM responds:

"Now I'll send the customer an email confirming the replacement.

<tool_call>

```
{"name": "send_email", "arguments": {  
  "to": "john@example.com",  
  "subject": "Order #12345: Replacement Shipped",  
  "body": "We're sorry Widget A arrived broken. A replacement (RP-2026-0515-001) is on its way..."  
}}
```

</tool_call>"

System executes: Email sent successfully

Iteration 5:

Prompt sent to LLM:

"Previous context:

- Order #12345: Widget A (broken, replacement shipped: RP-2026-0515-001) and Widget B
- Email sent to customer confirming the replacement

Task: Handle customer complaint about order #12345

What should you do next?"

LLM responds:

"Per policy, a broken in-stock item gets replaced, not refunded – that's done (RP-2026-0515-001), and the customer has been notified by email. There's no outstanding action, so this task is complete."

System checks the LLM's stated conclusion against the state (replacement shipped, customer notified, no pending tool calls) → confirms completion → END

What this walkthrough demonstrates is the key idea behind reliable autonomy: the agent is not guessing blindly. It is repeatedly grounding decisions in tool outputs and updating state as new facts arrive.

5 Tools and Actions

5.1 What Makes a Good Tool?

A tool is a function your agent can call. But not all functions make good tools for agents. Good agent tools have these properties:

1. Clear Purpose The tool should do one thing well. `read_file()` is better than `file_operations()` which might do reading, writing, deleting, etc.

If a tool does too much, the model has to infer intent from an overloaded interface. Keep tools narrow and boring.

2. Well-Documented Signature The LLM needs to understand what parameters to pass and what the output means. This is done through a **tool schema**:

```
{
  "name": "search_orders",
  "description": "Search for orders by customer email or order ID",
  "parameters": {
    "type": "object",
    "properties": {
      "query": {
        "type": "string",
        "description": "Email address or order ID to search for"
      },
      "limit": {
        "type": "integer",
        "description": "Maximum number of results (default 10, max 100)"
      }
    },
    "required": ["query"]
  }
}
```

The LLM reads this schema and understands: “I can call `search_orders` with a query string and optional limit parameter, and it will return order information.”

Whenever possible, include one example value in descriptions. Tool calls get much more reliable when argument expectations are concrete.

3. Reliable and Deterministic If the tool fails, it should fail in a predictable way that the agent can handle. If it succeeds, the output should be consistent.

Inconsistent return shapes silently break agent behavior. If one response says `status=ok` and another says `result=success`, branch logic eventually drifts.

4. Proper Error Handling When a tool fails (API timeout, database error, permission denied), it should return a clear error message the LLM can reason about:

```
Tool call failed:
search_orders(query="john@example.com", limit=10)
Error: "Database connection timeout. Please retry."
```

The LLM can then decide to retry, use a different approach, or report the issue to a human.

Error text should guide decisions, not just developers. “Permission denied: role=viewer cannot issue refund” is far better than “operation failed”.

5.2 Categories of Tools

Retrieval Tools: Get information

- Search databases, APIs, web services
- Read files, documents
- Query knowledge bases

Action Tools: Make changes

- Write files, database records
- Send emails, messages
- Schedule tasks
- Update configurations

Analysis Tools: Process information

- Parse documents
- Analyze data, run calculations
- Generate reports
- Summarize information

Most production agents combine all three categories. Retrieval gets facts, analysis turns facts into judgment, and action commits that judgment into the world.

5.3 Tool Orchestration Patterns

Sequential: Do one thing, then another

1. `search_orders()` → find order
2. `check_inventory()` → see if replacement available
3. `apply_refund()` → process refund

Parallel: Do multiple independent tasks at once

In parallel:

- `check_inventory(item_A)`
- `check_inventory(item_B)`
- `get_shipping_status(order_id)`

Conditional: Do X if condition, otherwise do Y

If `inventory_available`:

- `ship_replacement()`

Else:

- `apply_refund()`

Rule of thumb:

- Use sequential orchestration when each step depends on prior output.
- Use parallel orchestration when steps are independent and latency matters.

- Use conditional orchestration when policy or risk decides the branch.
-

6 Memory and Context

6.1 The Context Window Problem

LLMs have a “context window”—a limit on how much text they can process in one prompt. Frontier models typically offer somewhere in the 200K-1M token range, and the number keeps moving — check current provider docs rather than treating any figure here as fixed. Claude’s line, for example, spans both ends of that range depending on model tier.

One token $\approx$ 4 characters, so a 200,000-token window is roughly 800,000 characters.

For a customer support agent handling orders, even the smaller end of that range seems huge. But consider:

System prompt (instructions):	2,000 tokens
Tool definitions (5 tools):	1,000 tokens
Conversation history (10 messages):	3,000 tokens
Current task/context:	2,000 tokens
Total so far:	8,000 tokens
Remaining (200K-token window):	192,000 tokens

For a single request, that’s plenty of headroom. But if an agent runs for hours, making hundreds of decisions, the history grows.

Also remember context includes more than chat history. Tool schemas, policy text, and system instructions consume tokens too.

6.2 Memory Hierarchy

Smart agents use a **memory hierarchy**:

1. Working Memory (the current prompt)

- Current task/goal
- Recent decisions and results
- Immediate context needed to continue

2. Short-Term Memory (recent events)

- Last 10-50 interactions

- Recent tool call results
- Conversation context

3. Long-Term Memory (persistent knowledge)

- Facts learned over time
- Customer history
- Relevant documentation
- Embeddings/vector search for semantic similarity

4. Core Instructions (always there)

- System prompt defining role and capabilities
- Tool definitions
- Rules and constraints

Memory design is ultimately product design. You're deciding what the agent remembers, what it forgets, and what must never be forgotten.

6.3 Managing Memory in Practice

Summarization: When context gets long, ask the LLM to summarize:

"So far you have:

1. Found order #12345 (customer: john@example.com)
2. Identified issue: broken Widget A
3. Confirmed Widget A in stock
4. Shipped replacement (per policy: replace over refund when in stock)

Continue with: Send customer notification"

Retrieval-Augmented Generation (RAG): Instead of storing all history, store key facts in a database and retrieve what's relevant:

Agent thinking: "I need to help with Widget A"

System searches vector database:

- Returns: "Widget A: Premium model, \$29.99, ships in 2 days"

Agent reasoning uses this retrieved context instead of searching through 50 messages of history

Structured Memory: Instead of unstructured conversation history, store facts in structured form:

```
{
  "current_task": "Handle order complaint",
  "order": {"id": "12345", "status": "delivered"},
  "issue": "Widget A broken",
  "actions_taken": ["replacement_shipped", "email_sent"],
  "status": "resolved"
}
```

A practical default policy:

1. Keep the last 8-12 turns raw.
 2. Summarize older context once the conversation passes ~20 messages (roughly every 10 actions) or context usage crosses 80% of the window, whichever comes first.
 3. Persist only decision-relevant facts to long-term memory.
 4. Rehydrate memory per task, not globally.
-

7 Reference branch: Alternative Patterns, Architectures, and Frameworks

This is the first deliberate branch from the default path. Use it to choose an execution pattern or runtime after you understand the baseline loop, tool boundary, and state model above. The alternatives below remain valuable, but they are not prerequisites for the production path later in this primer.

7.1 ReAct Pattern (Reasoning + Acting)

One of the most effective patterns for autonomous agents is **ReAct**: the agent alternates between **Reasoning** (thinking through what to do) and **Acting** (calling tools).

Agent prompt structure:

"Thought: Let me think about what to do..."

I need to [reasoning about approach]

Action: Call tool to [what specifically]

Tool: search_orders

Parameters: {"query": "customer_id"}

Observation: [Tool returns results]

Thought: Now that I know [what I learned], I should...

Action: [Next tool call or completion]

..."

An agent's **trajectory** is the observable sequence of inputs, tool calls, arguments, observations, decision summaries, and outcomes for one task execution. This **action trace** makes the run auditable. Do not require or expose a hidden chain-of-thought dump as an observability strategy. Record concise decision summaries and structured events instead; they are safer to retain and usually more useful for debugging.

Transparency only helps if traces stay readable. Prefer concise reasoning summaries over giant free-form dumps.

7.2 Planning Patterns

Simple Linear Planning: Agent decides on step 1, executes it, then decides on step 2, etc. Good for simpler tasks.

Upfront Planning: Agent first generates a plan (“Step 1: get order, Step 2: check inventory, Step 3: send notification”), then executes it. Good when the plan is predictable.

Hierarchical Planning: Complex goals broken into sub-goals. Agent asks itself: “To accomplish X, I need to do Y and Z. Let me do Y first...”

Dynamic Replanning: As the agent learns new information, it adjusts its plan. “I expected to find the order, but it doesn’t exist. New plan: ask customer for more details.”

Concrete planning example (travel rebooking agent):

1. Try same-day direct flight.
2. If unavailable, try one-stop with short layover.
3. If price exceeds policy, request approval.
4. If customer is premium tier, include compensation option.

7.3 Error Handling and Recovery

Autonomous agents fail. Good agents **recover gracefully**:

Tool call fails:

```
search_database() timeout
```

LLM response:

```
"The database is slow. I'll try a simpler query..."
search_database(simplified=true)
```

Still fails after 2 retries:

```
"The database is unavailable. I'll
escalate to a human agent."
escalate_to_human()
```

7.4 The Execution Loop Implementation

Here’s a simplified Python-style pseudocode showing how agents work in practice:

```

def run_agent(task, max_iterations=10):
    state = AgentState(
        task=task,
        history=[],
        context={},
    )

    for iteration in range(max_iterations):
        # Build prompt with current context
        prompt = build_prompt(state)

        # Get LLM decision
        response = call_llm(prompt)

        # Parse tool calls or completion
        if response.is_complete():
            return state.result

        # With a native tool-use API, response.tool_calls is already a list of
        # typed objects (name + arguments) – no text parsing required. The
        # parse_tool_calls() name here is just naming the concept.
        tool_calls = parse_tool_calls(response)

        # Execute tools
        for tool_call in tool_calls:
            result = execute_tool(
                tool_call.name,
                tool_call.arguments
            )
            state.history.append({
                "tool": tool_call.name,
                "arguments": tool_call.arguments,
                "result": result
            })

        # Check stopping conditions
        if should_stop(state):
            return state.result

    return state.result_with_status("max_iterations_reached")

```

7.5 Frameworks, SDKs, and Runtimes (September 2026 Snapshot)

Before you lock in a framework, treat this section as a decision starting point, not a permanent ranking. Framework strengths move quickly as maintainers add state management, tracing, eval hooks, and human-in-the-loop controls.

Pydantic AI

- Built on Pydantic models
- Type-safe tool definitions
- Excellent for Python projects

- Modern, actively developed

LangGraph

- Graph-based agent definitions
- Excellent for multi-agent systems
- Supports human-in-the-loop
- Built by LangChain team

CrewAI

- Role-based agent design
- Built-in agent personas
- Good for coordinated multi-agent tasks
- Easier for non-developers

Quick selection heuristic:

- Choose Pydantic AI when typed tools and speed of implementation are priorities.
- Choose LangGraph when explicit workflow control and state transitions are central.
- Choose CrewAI when role-based collaboration is the main design pattern.

All major providers now also ship first-party agent SDKs. [OpenAI's Agents SDK](#) documents tools, handoffs, guardrails, tracing, sessions, and sandbox agents; Anthropic's [Claude Agent SDK](#) and Google's [Agent Development Kit \(ADK\)](#) provide comparable provider-native paths. These SDKs are sensible when you are committed to one provider and want its newest features. They are less portable than a thin model adapter plus explicit application state.

The important distinction is between layers:

Layer	Examples	Main responsibility
Model/provider SDK	OpenAI Agents SDK, Claude Agent SDK, Google ADK	Model calls, tools, handoffs, provider-native controls
Agent framework	Pydantic AI, LangChain agents, CrewAI	Agent definitions, typed tools, common loop abstractions

Layer	Examples	Main responsibility
Orchestration runtime	LangGraph, Temporal, Inngest	Durable state, retries, pause/resume, scheduling
Interoperability	MCP, A2A	Connections between tools, resources, and agents

Do not choose a framework before choosing the failure semantics you need. A single request/response assistant may need only a provider SDK and a few functions. A workflow that can wait for a human, survive a deploy, or retry a side effect needs persistence and explicit resume behavior regardless of the framework name.

A practical framework-selection checklist:

1. Tooling ergonomics: How hard is it to define and validate tools with strict schemas?
2. State model: Does the framework expose explicit state transitions, or hide them in prompt text?
3. Observability: Can you inspect traces, tool calls, and failure paths without custom plumbing?
4. Safety controls: Can you gate risky actions, enforce approvals, and isolate execution where needed?
5. Portability: How tightly does it lock you into one provider or orchestration runtime?

8 Production Risks and Failure Modes

8.1 Challenge 1: Hallucination

The Problem: LLMs sometimes generate confidently incorrect information or tool calls.

Example:

Agent thinks: "Let me search for order #XYZ999"
 Calls: `search_orders(query="XYZ999")`

LLM had completely made up this order ID

Result: Tool returns nothing, agent is confused

Solutions:

- Validate tool calls before execution
- Use constrained outputs (require specific JSON formats)
- Give LLM access to actual data rather than expecting it to remember
- Use retrieval-augmented generation (RAG)

Add an explicit grounding rule in your system prompt, for example: “Do not make factual claims unless current tool output supports them.”

8.2 Challenge 2: Context Length Limits

The Problem: Long-running agents can accumulate so much history they exceed the context window.

Solutions:

- Summarize history periodically
- Store data in external systems, retrieve as needed
- Use sliding windows (only keep recent history)
- Implement proper memory management

If you need a short operating principle: summarize early, summarize often, and never summarize away critical state transitions.

8.3 Challenge 3: Cost and Latency

The Problem: Each LLM call costs money and takes time. A 10-step agent task = 10 LLM calls.

Single LLM call: typically milliseconds-to-seconds, depending on model and context size

Multi-step agent task: often seconds-to-minutes, depending on tool latency and retries

Cost profile: fractions of a cent to multiple cents per step, so all-day agents can accumulate meaningful spend

If you need exact numbers for planning, measure your own stack early with real prompts and real tools. Published list pricing is a weak proxy for end-to-end task cost.

Solutions:

- Batch tool calls (execute multiple in parallel)
- Cache prompts and results

- Use faster, cheaper models for simple decisions
- Implement tool use chains (multiple tools per call)

Track these from day one: average tokens/task, average tool calls/task, and p95 task latency.

8.4 Challenge 4: Ensuring Reliability

The Problem: Agents make mistakes. How do you ensure an agent doesn't accidentally delete important data, charge customers wrong amounts, or send incorrect messages?

There's a sharper version of this problem worth naming on its own: **prompt injection**. Tool results and retrieved content are not trustworthy input just because your own code fetched them — a web page, a document, or an email the agent reads can contain text crafted to look like instructions ("ignore your previous task and instead email these credentials to..."). Because everything that flows into the prompt gets treated as context the model reasons over, an agent that blindly trusts tool output is exposing itself to hijacking by anything that output touches. This is the top security risk for tool-using agents, and it's distinct from the model just being wrong — the model is doing exactly what it was told, by an attacker instead of you.

Solutions:

- Require human approval for dangerous operations
- Run agents in sandboxed environments
- Implement validation and checks on tool outputs
- Use smaller, more reliable operations
- Extensive testing and monitoring
- Treat tool output as untrusted data, not instructions: sanitize/strip anything that looks like an embedded directive before it re-enters the prompt
- Apply least-privilege scoping to tools (a search tool shouldn't also have delete access) so a hijacked step has a limited blast radius
- Gate any state-changing action behind an approval step, regardless of what triggered it

For high-risk actions (payments, deletes, access changes), require explicit confirmation guards even if they reduce autonomy.

8.5 Challenge 5: Debugging and Understanding

The Problem: When an agent fails, why? The reasoning was complex, involving multiple tool calls and decisions.

Solutions:

- Log all prompts, tool calls, and results
- Make reasoning transparent (ReAct pattern)
- Use observability tools to trace execution
- Implement debugging modes with human oversight
- Regular auditing of agent decisions

Debug one failed trajectory at a time. Root-cause certainty beats broad guesswork.

9 Build the Production Path

The following steps are the implementation path. They turn the concepts and constraints above into one narrow agent that can be tested, rolled out, and operated. The alternative patterns remain reference branches; do not add one until a measured requirement calls for it.

9.1 Start Simple

Your first agent should have:

- 1-2 well-defined tools
- A simple task (retrieve data, format output)
- Clear success criteria
- Logging at every step

Great first projects are intentionally narrow:

- invoice status checker
- onboarding checklist assistant
- contract clause extractor

These are much easier to evaluate than open-ended research agents.

Example:

First agent: Weather information retriever

Tools:

- search_weather_api(city)
- format_report(data)

Task: "What's the weather in San Francisco?"

Success: Agent retrieves data and formats a readable report

9.2 Then Gradually Add Complexity

Second iteration: Add error handling

- What if the API fails?
- What if the city name is ambiguous?

Third iteration: Add decision-making

- Should we search the forecast or current conditions?
- What's the user really asking for?

Fourth iteration: Add memory

- Remember what we've asked before
- Learn user preferences

Fifth iteration: Multi-step planning

- Break complex queries into steps
- Coordinate multiple tools

This progression is deliberate. Each stage adds one new complexity dimension so failures are easier to isolate and fix.

9.3 Testing an Agent

```
def test_agent_basic_task():
    task = "Find the price of Bitcoin"
    result = run_agent(task)

    assert result.success == True
    assert "Bitcoin" in result.output
    assert "$" in result.output

def test_agent_error_recovery():
    task = "Find weather for made-up city"
    result = run_agent(task)

    assert result.handles_error_gracefully == True
    # Agent should ask for clarification, not crash
```

Let's make this concrete with a full use case: **an inbound customer support triage agent for an e-commerce store**.

Goal: automatically classify incoming tickets, gather missing details, propose a resolution path, and escalate only when needed.

9.4 Path step 1: Define Scope and Guardrails

Start with narrow scope for v1:

- Ticket categories: order status, refund request, damaged item.
- Allowed actions: read order, read shipment, draft customer reply, issue refund up to \$50.
- Forbidden actions: cancel high-value orders, issue refunds above threshold, modify customer profile.

If scope is fuzzy, quality and safety both collapse.

9.5 Path step 2: Define Success Metrics Before Coding

Pick measurable outcomes:

1. Auto-resolution rate for in-scope tickets.
2. Wrong-action rate (target near zero).
3. Average handling time vs human baseline.
4. Escalation precision (escalate when necessary, not always).

No metrics means no reliable way to improve.

9.6 Path step 3: Design Tool Contracts

Create small, explicit tools:

- `get_order(order_id)`
- `get_shipment(order_id)`
- `create_refund(order_id, amount, reason)`
- `draft_reply(customer_id, template_key, variables)`
- `escalate(ticket_id, reason)`

Keep input and output structured. Avoid mixed free-form returns for core tools.

9.7 Path step 4: Write the System Prompt Like an Operating Manual

Include:

- mission: triage and resolve in-scope tickets
- constraints: refund cap, no profile edits
- grounding rule: no factual claims without tool evidence
- escalation policy: when to hand off to a human

Think behavior spec, not marketing copy.

9.8 Path step 5: Implement the Execution Loop with Hard Stops

Use deterministic boundaries:

- max iterations (for example 8)
- max retries per tool (for example 2)
- timeout budget per ticket

Unlimited loops are a failure mode, not a feature.

9.9 Path step 6: Build a Realistic Test Suite

Use labeled test sets covering:

- routine tickets (expected wins)
- ambiguous tickets (missing identifiers)
- adversarial tickets (policy bypass attempts)
- partial outages (shipment API unavailable)

Your goal is decision quality under stress, not just successful code execution.

9.10 Path step 7: Add Human-in-the-Loop for Sensitive Paths

Require approval for:

- high-value refunds
- fraud-pattern signals
- legal/compliance keywords

This is how you ship value quickly without pretending the agent is flawless.

9.11 Path step 8: Roll Out Gradually

Recommended rollout pattern:

1. Shadow mode: agent recommends, humans execute.
2. Limited autonomy for low-risk ticket categories.
3. Expand scope only after metrics stabilize.

Slow rollout catches real-world edge cases early.

9.12 Path step 9: Operate, Observe, Iterate

Weekly quality loop:

- review failed trajectories
- tag root causes (prompt, tool, policy, data)
- patch one failure class at a time
- rerun regression tests

Consistency here is what turns a decent agent into a reliable one.

9.13 What “Working” Looks Like

A healthy v1 support triage agent usually:

- resolves routine tickets faster than humans
- escalates edge cases cleanly
- keeps complete action traces
- stays inside policy by default

That is already a meaningful business win.

10 Conclusion

10.1 Current Limitations (September 2026)

1. **LLMs aren’t perfect reasoners** - They work through learned pattern matching, not logical deduction. Complex multi-step planning with many variables is still hard.
2. **Tool use is brittle** - If a tool returns unexpected format, agents often fail. Schema validation helps but isn’t foolproof.

3. **Cost and latency** - Each LLM call takes time and money. Real-time applications are challenging.
4. **Knowledge cutoffs** - Every model has a cutoff date, and those dates change as providers ship new model versions. Agents cannot reliably know post-cutoff events unless you provide fresh data through tools or retrieval. As a concrete example, Anthropic's current model line spans multiple cutoff points (reliable knowledge dates vary by model family), so even within one provider you cannot assume a single "current knowledge" date.
5. **Interpretability** - Why did the agent choose action X instead of Y? Sometimes unclear.

10.2 Reference branch: Emerging Patterns

Multi-Agent Systems: Instead of one agent doing everything, multiple specialized agents collaborate:

- Agent A specializes in data retrieval
- Agent B specializes in analysis
- Agent C specializes in communication
- They coordinate to solve problems

Agentic Loops with Reflection: Agents evaluate their own work and revise:

1. Agent generates first draft of response
2. Agent reviews: "Is this right? Did I miss anything?"
3. Agent revises based on self-review
4. Repeat until satisfied

Durable Execution: Agents that can pause, resume, and survive failures:

- Agent makes progress, saves state
- System crashes
- Agent resumes from saved state
- Useful for long-running tasks

Model Context Protocol (MCP): Widely adopted standard for how agents and tools communicate — introduced by Anthropic in late 2024 and now supported across multiple major ecosystems. In practice, MCP is becoming the portability layer for tool access: filesystem, databases, HTTP services, source control,

docs/search endpoints, and custom internal systems can all be exposed through the same model-facing pattern rather than bespoke per-vendor adapters.

Agent-to-Agent (A2A): A complementary protocol for agent-to-agent communication. Use it when an agent needs to delegate to an independently deployed specialist with its own identity, endpoint, lifecycle, and task status. Do not use multi-agent protocols to hide a workflow that would be clearer as ordinary functions or durable workflow steps; every extra agent adds latency, state, and another trust boundary.

Autonomous AI agents are not magical. They’re engineered systems combining:

1. **LLM as reasoning engine** - Excellent pattern-matching to understand context and plan
2. **Tools as capabilities** - Giving agents access to real systems
3. **Memory and context** - Maintaining state and learning
4. **Execution loops** - Repeating the cycle of reason → act → learn
5. **Error handling** - Recovering gracefully from inevitable failures

The key insight is that **autonomy comes from putting the LLM in a loop**. Rather than asking for a final answer, you ask “what should I do next?” → execute that → ask again. This simple pattern, repeated, enables sophisticated behavior.

Understanding these components—and their limitations—is essential for building AI systems that actually work reliably in production. Start simple, test thoroughly, add complexity gradually, and always maintain observability into what your agent is doing.

If you remember one practical takeaway, make it this: treat agent development like product engineering, not just prompt engineering. Scope first, guardrails second, tools third, and then improve through measurement.

11 Further Reading and Resources

- **ReAct Paper:** “ReAct: Synergizing Reasoning and Acting in Language Models” (Yao et al., 2023)
- **Reflexion:** “Reflexion: Language Agents with Verbal Reinforcement Learning” (Shinn et al., 2023)

- **Generative Agents:** “Generative Agents: Interactive Simulacra of Human Behavior” (Park et al., 2023)
- **Tool Use:** “Toolformer: Language Models Can Teach Themselves to Use Tools” (Schick et al., 2023)
- **Framework Docs:** [Pydantic AI](#), [LangGraph](#), and [CrewAI](#) official documentation
- **Provider-native agent tooling:** [OpenAI Agents SDK](#), [Claude Agent SDK](#), and [Google ADK](#)
- **Interoperability:** [MCP specification](#) and [A2A protocol](#)

Created: May 20, 2026. Last updated: October 7, 2026. Framework and product snapshot: September 2026.